



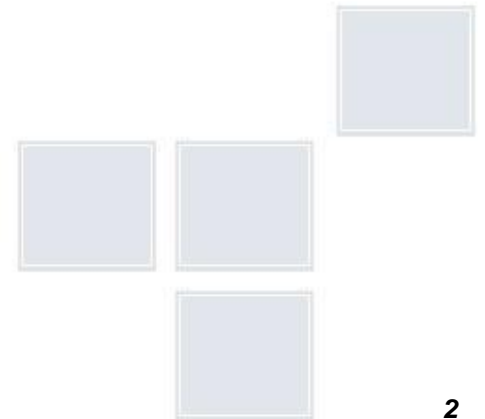
# Process

COMP3230B Workshop 2



# Learning Outcomes

- To learn the basics of multiprocessing programming
- To understand the flow of Assignment 1



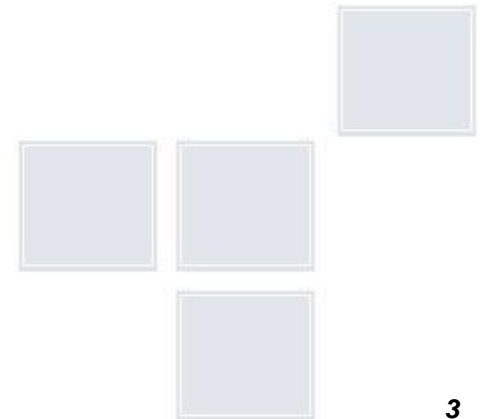


# Workshop Setup

■ **Time: 13:30 – 15:20**

■ **In-Class Preparation**

- **Boot the Linux system or VM in your laptop, or;**
- **Connect CSVPN and log in X2Go**



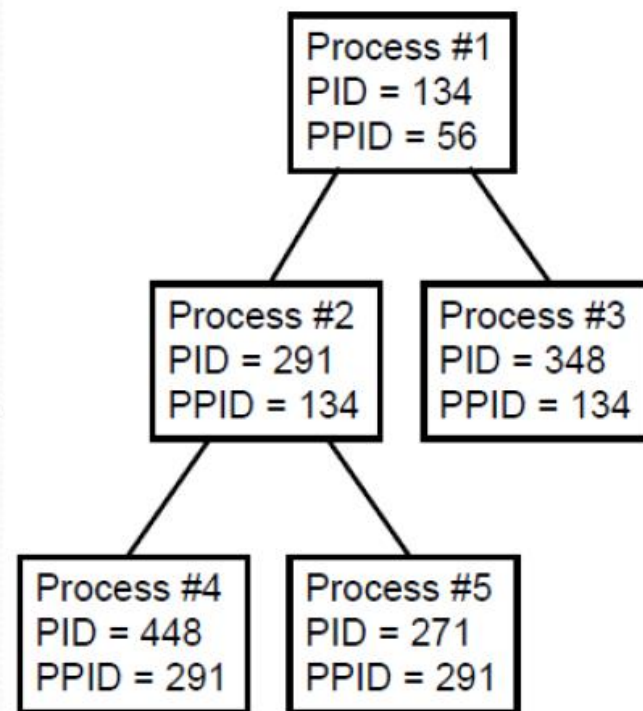


# Multiprocess Programming



# Linux Process

- Each process has a unique Process ID (**PID**) that allows the kernel to identify it.
- Each process may start an unlimited number of other processes called child processes.
- Each process has only one parent process.
- Each process must have been started by an existing process called a parent process.
- Each process has a Parent Process ID (**PPID**) which identified the process that started it.
- PID is generated randomly from free entries in a process table used by the Linux kernel.



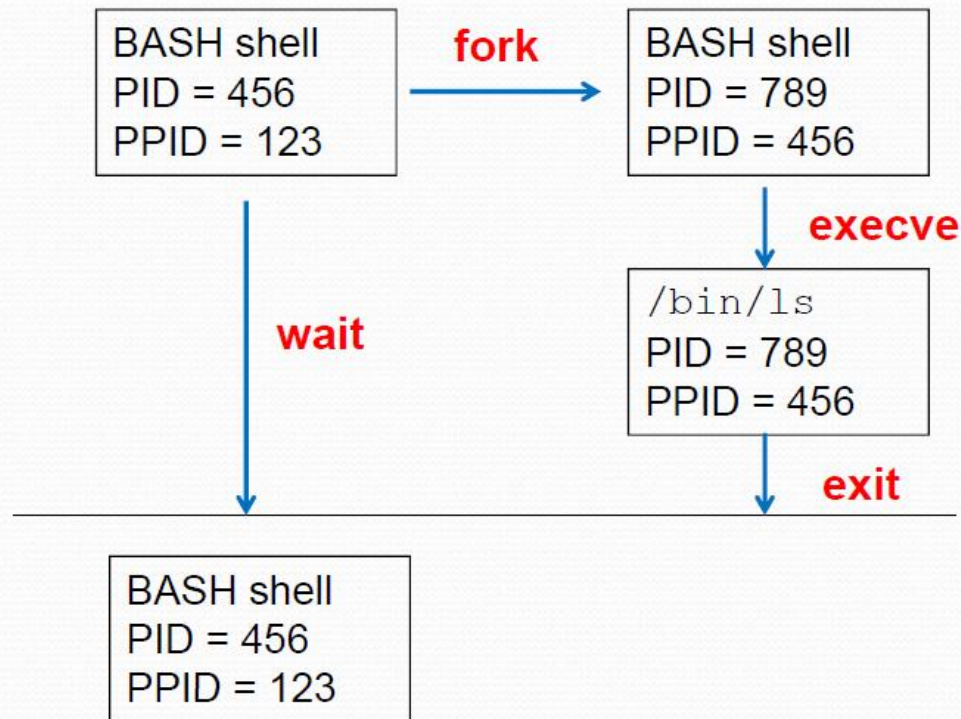
**Relationship between  
Parent and Child Processes**





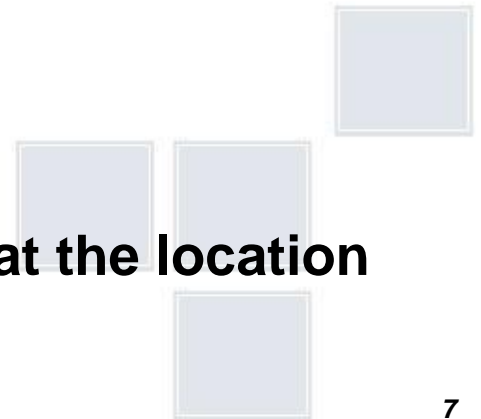
# Process Creation and Termination

- How to start a process?
  - E.g. executing `ls` through a shell





- **System call to create a new process (child) by duplicating the caller process.**
- **Format: `pid_t fork();`**
- **Each process receives different return value from `fork()`.**
  - **In the caller process, the return value is the PID of the child process.**
  - **In the child process, the return value is 0.**
- **Two processes will run "concurrently"**
  - **In the child process, the execution starts at the location after `fork()`.**



E.g. fork.c

```
#include <stdio.h>
#include <sys/types.h>

int main () {
    pid_t pid=fork();

    if (pid < 0) {
        printf("can't create process\n");
    } else if (pid == 0) {
        printf("I am child\n");
    } else {
        printf("I am parent\n");
    }
    return 0;
}
```

Output:

```
I am parent
I am child
```



## Parent Process

a new process  
is created

## Child Process

```
→ pid_t pid = fork();  
→ if (pid < 0) {  
    printf("can't create  
    process\n");  
}  
→ else if (pid == 0) {  
    printf("I am child\n");  
}  
→ else {  
    printf("I am parent\n");  
}  
→ ...
```

```
→ pid_t pid = fork();  
→ if (pid < 0) {  
    printf("can't create  
    process\n");  
}  
→ else if (pid == 0) {  
    printf("I am child\n");  
}  
→ else {  
    printf("I am parent\n");  
}  
→ ...
```



# getpid() / getppid()

- System call
- Function: get pid and ppid of the current process
- Format:
  - `pid_t getpid();`
  - `pid_t getppid();`
- Returned type: `pid_t`
  - Treat it as integer

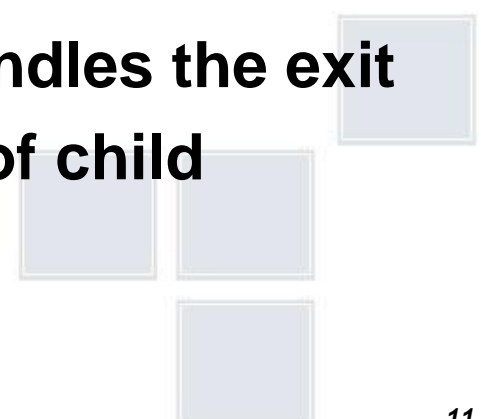
```
pid_t my_pid, parent_pid;

my_pid = getpid();
parent_pid = getppid();
printf("My pid is %d\n", my_pid);
printf("My parent's pid is %d\n", parent_pid);
```



# Example: Zombie

- Download zombie.c
- Compile and run
- In a new terminal:
  - `$ ps aux | grep <user name>`
- Zombie is not actually existing any more. Its record is kept until parent reads it.
  - Zombie processes only consume some space in process table. **The process with Z+ status**
- The child finishes, but the parent is in deep sleep
- The entry in the process table is kept until the parent wakes up and handles the exit message of child







# exec family

- exec is actually family: 6 functions.

- Only *execve* is *system call*.
- Other five are *library call* that calls *execve*.
- Example:

```
execlp("/bin/ls", "ls", NULL);
```

- As soon as a program (process) call **exec**, most of its states will be **replaced** by the new command-process, except a few information, such as PID.

```
int execl(const char *path, const char *arg, ...);  
int execlp(const char *file, const char *arg, ...);  
int execle(const char *path, const char *arg, ..., char *const envp[]);  
int execv(const char *path, char *const argv[]);  
int execvp(const char *file, char *const argv[]);  
int execve(const char *path, char *const argv[], char *const envp[]);
```

Text segment

Data segment

stack

.....

*We will try  
this in  
Exercise*

- Function: used by parent to wait for a child process to terminate.
  - Can't wait for a process which is not started by the caller.
- Format:
  - `pid_t wait(int *status);`
  - `pid_t waitpid(pid_t pid, int *status, int options);`
    - *pid\_t* means process id.
    - *status* is used to store the child's status information, such as whether the child terminated normally.
    - *options*: WNOHANG /WUNTRACED
- Returned value: process id of child, or -1 for error.
  - What error? Such as:
    - What if the process specified in pid does not exist.
    - What if the options argument was invalid.
  - Parent (the caller) wait until child exit.
  - If you don't care the status, you can use *NULL* for simplicity
    - E.g. `waitpid(123, NULL, 0);`





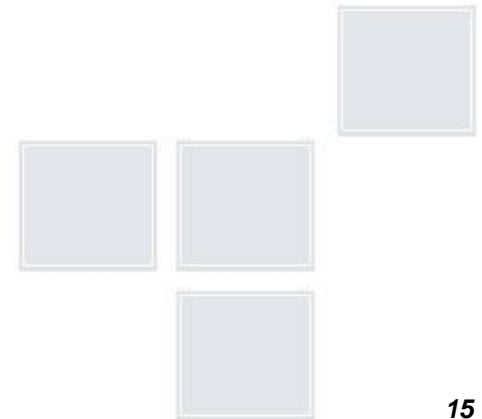
# What should be included as header??

- **To view manual pages, issue commands in terminal**
  - **man <optional section number> <manual page>**
  - **e.g. man fork, man 3 printf**
- **There may be manual pages with the same name but different section numbers, e.g. printf(1), printf(3)**
  - **1 for commands, 2 for system calls, 3 for library calls, ...**
- **To search for pages, use -k option.**
  - **man -k sort**





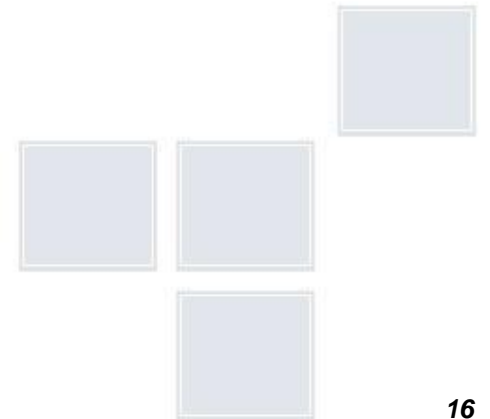
- **Similar to navigation in Vi editor**
  - **j (or down-arrow) for one line down, k (or up-arrow) for one line up**
  - **Ctrl-v for next page, Ctrl-b for previous page**
  - **xxG for xx-th line, where xx is integer.**
  - **/abc for searching "abc", then n for next match, and N for previous match**
  - **h for help, q for quit**





# Keywords in Manual Pages

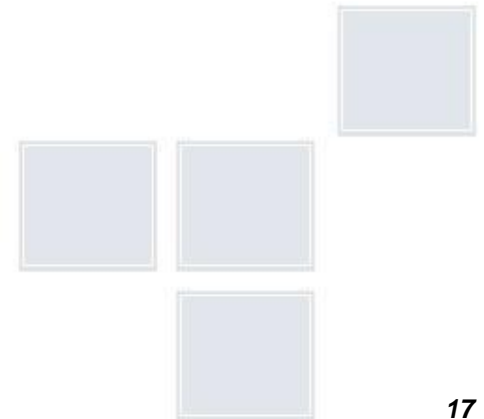
- **Synopsis**
  - What headers should I include?
  - What are the parameters?
- **Description**
  - What does the function do?
- **Return Values, Errors**
  - How do I know the execution result?
- **See Also**
  - What are the related functions?
- **Standards**
  - Is it still useful on another OS?





# Example: Make Use of Linux Command

- In **fork.c** (Task 1 of Exercise)
  - Let child exec the `ls` command (As shown in Page 12)
  - Let parent wait for child
- In this way, programmers can use Linux commands as part of their programs.
  - e.g. Use `sort` to sort an input -> wait -> proceed to read the sorted file





# Assignment 1



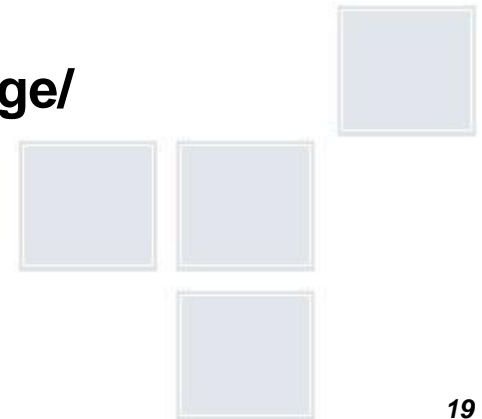


## ■ Monitor the run-time memory usage

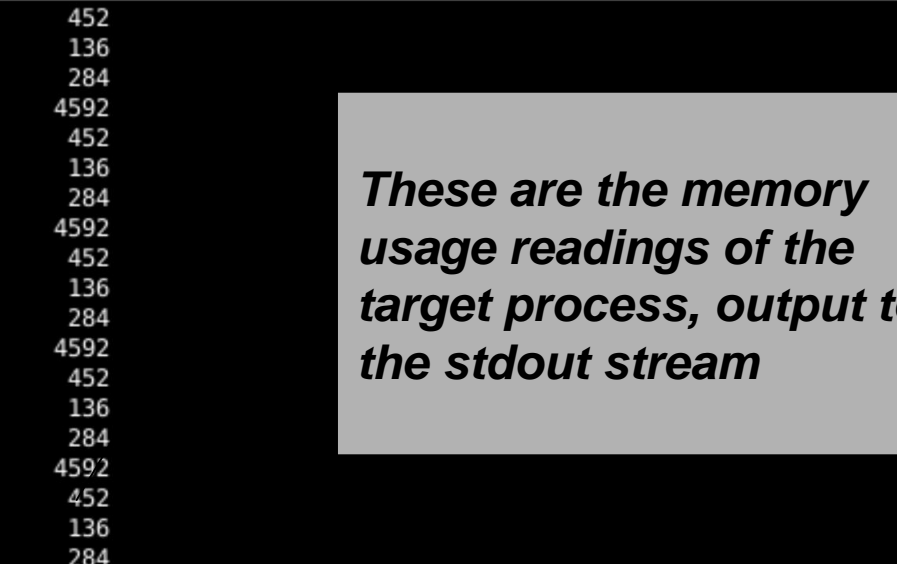
- To access the total memory information about a process on Linux, we use the /proc virtual file system. Within it there is a directory full of information for each active process id (pid). By reading `/proc/(pid)/status` we can obtain information about memory.

## ■ Credit to Lockless

- [http://locklessinc.com/articles/memory\\_usage/](http://locklessinc.com/articles/memory_usage/)





- 
- The screenshot shows a terminal window with a blue title bar labeled "Terminal". The menu bar includes "File", "Edit", "View", "Terminal", "Tabs", and "Help". The terminal output displays a repeating pattern of memory usage readings for a target process, output to the stdout stream. The readings are as follows:
- | Process ID | Memory Usage |
|------------|--------------|
| 1:         | 452          |
| 2:         | 136          |
| 3:         | 284          |
| 0:         | 4592         |
| 1:         | 452          |
| 2:         | 136          |
| 3:         | 284          |
| 0:         | 4592         |
| 1:         | 452          |
| 2:         | 136          |
| 3:         | 284          |
| 0:         | 4592         |
| 1:         | 452          |
| 2:         | 136          |
| 3:         | 284          |
| 0:         | 4592         |
| 1:         | 452          |
| 2:         | 136          |
| 3:         | 284          |
| 0:         | 4592         |
| 1:         | 452          |
| 2:         | 136          |
| 3:         | 284          |
| 0:         | 4592         |
- These are the memory usage readings of the target process, output to the stdout stream*



- A perl script that drives gnuplot

- Read labeled streams from stdout

*Pass stdout of process 1 to  
stdin of process 2*

- Usage

- `./tmem firefox | perl driveGnuPlots.pl 4 400 400 400 400`  
`vmsize vmdata vmstk vmrss`

- Before you try, close all the Firefox windows!

- Credit to ttsiodras

- <http://users.softlab.ntua.gr/~ttsiod/gnuplotStreaming.html>



# Running on own machine

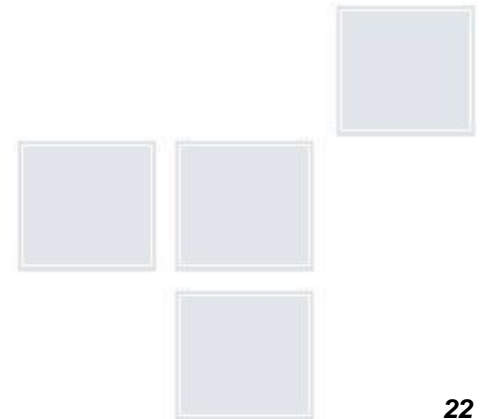
- **Install gnuplot and gnuplot-x11**

- **Ubuntu**

- `sudo apt-get install gnuplot gnuplot-x11`

- **Fedora**

- `sudo yum install gnuplot gnuplot-x11`









```
//Set the signal handler function as sig_chld
act.sa_handler = sig_chld;

/* We don't want to block any other signals */
sigemptyset(&act.sa_mask);

act.sa_flags = SA_NOCLDSTOP;

if (sigaction(SIGCHLD, &act, NULL) < 0)
{
    fprintf(stderr, "sigaction failed\n");
    return 1;
}
```

## ■ Set signal handler as sig\_chld

- Handler means a function that is called every time its corresponding signal arrives

## ■ SIGCHLD

- A signal that a child sends to parent when child finishes
- We may also send through the terminal



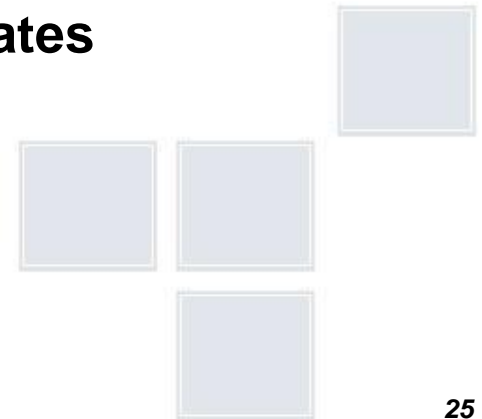
# sig\_chld()

```
pid = waitpid(-1, &status, WNOHANG);
if (pid < 0)
{
    fprintf(stderr, "waitpid failed\n");
    return;
}

/* pid blocked */
if(pid == block_pid)
    return;

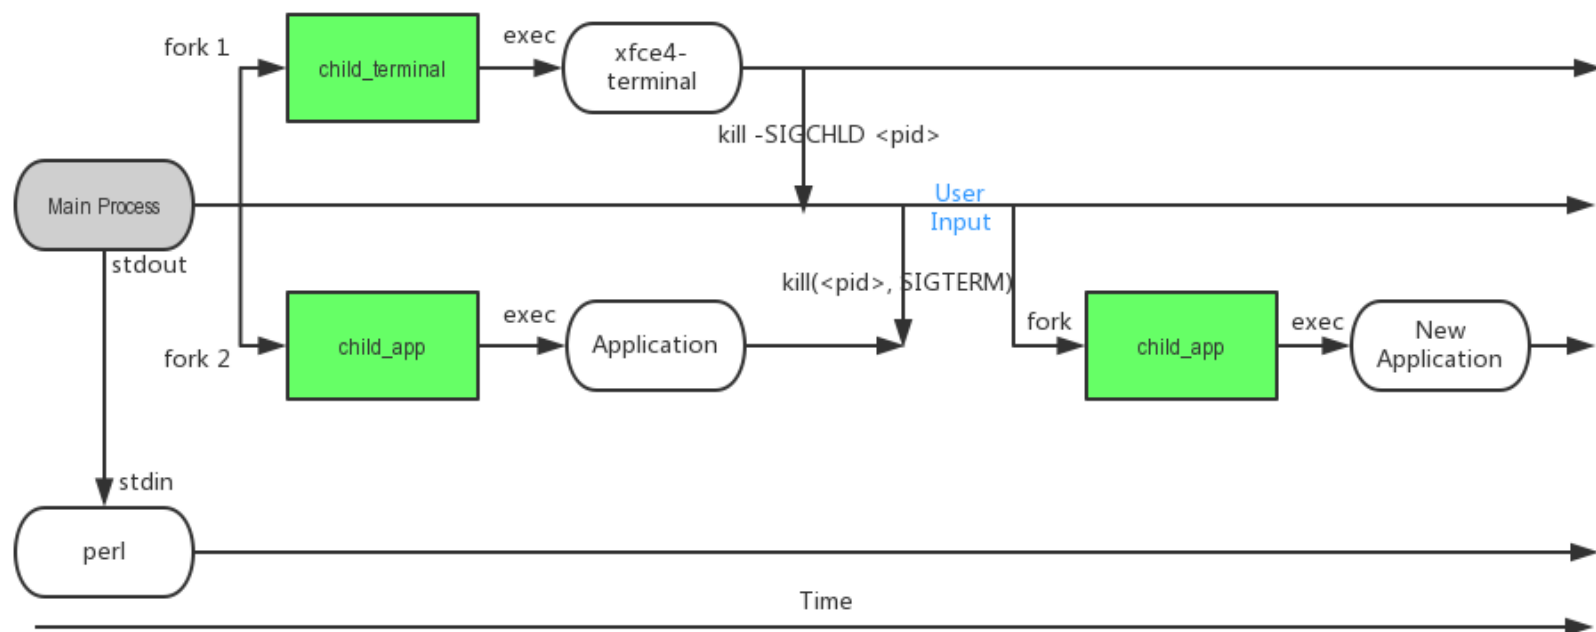
/* signal not from children */
if (pid != target_pid )
```

- **WNOHANG:** Enter handler immediately without looking at the sender pid
- Only terminate when direct child terminates





# At completion of assignment



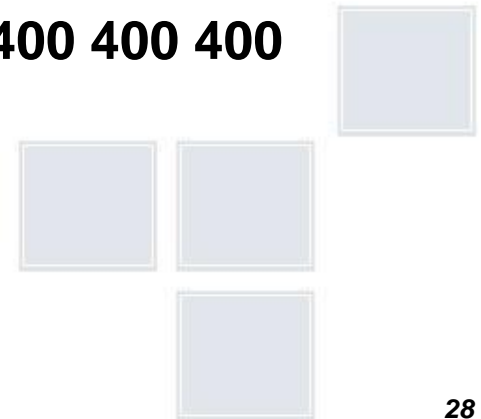


# Exercise



# Workshop Exercise 2

- **Task 1: Finish the task on Page 17 and submit `fork.c`**
  - **Exec the `ls` command. Not sort.**
- **Task 2: Submit **1** screenshot showing the running of `tmem` with real-time display**
  - We just want to make sure everyone can run the basic program
  - `./tmem firefox | perl driveGnuPlots.pl 4 400 400 400 400`  
`vmsize vmdata vmstk vmrss`







# Workshop Exercise 2

- **Modify example\_1.c to make sure the main process creates exactly 4 child processes**

```
// wrong example_1.c
#include <unistd.h>    // for fork()
#include <stdio.h>

int main() {
    fork();
    printf("Child 1 created\n");
    fork();
    printf("Child 2 created\n");
    fork();
    printf("Child 3 created\n");
    fork();
    printf("Child 4 created\n");
    return 0;
}
```

- **Submit the corrected `example_1.c` and 1 screenshot showing the correct output (output order does not matter)**